

Python Programming and Machine Learning for Economists (Jan/Feb 2022)

Michael E. Rose, PhD

Introduction

Who am I?

- Senior Research Fellow, Max Planck Institute for Innovation and Competition, PhD in Econ (University of Cape Town)
- Writing code since 8th grade
- Author of 3 open-source projects: `pybliometrics`, `sosia`, `scholarmetrics`
- Teaching experience:
 - *This course* @ Kiel Institute for the World Economy (ASP), University of Zurich, ifo Institute Munich, LMU Munich, Scheller College of Business at Georgia Tech, TU Munich
 - Risk Management Computing Skills [Matlab, SQL, Excel, VBA] @ University of Cape Town
- Michael.Ernst.Rose@gmail.com



Who are you?

- Name, Status
- Which languages, how long?
- Which operating system?
- Who is more in control, your computer or you?

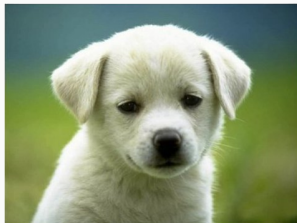
1. Empirical research using Python
2. Project management
3. Supervised Machine Learning
4. Unsupervised Machine Learning
5. Natural Language Processing

- Lecture in the morning, exercises in the afternoon
- Each exercise session starts with a Monty Python sketch
- 10 Minutes breaks after 50 Minutes of Teaching

Exercises (= mini projects)

👍 Difficulty increases as the course progresses

Data sets
in tutorials



Data sets in
the wild



👍 Your grades depend on the final exercises

Learning outcomes

- Programming part
 1. List some of the right basic tools for empirical research
 2. Use python independently
 3. Apply pandas, seaborn, sklearn
 4. Understand coding principles
 5. Use PyCharm
 6. Understand version control and use git
- Machine Learning
 1. Apply simple Neural Networks, clustering algorithms and Principal Component Analysis
 2. Interpret and evaluate any machine learning application
 3. Teach yourself how to apply machine learning algorithms we don't speak about

Required Readings

- 📖 Shapiro, J. and M. Gentzkow: “Code and Data for the Social Sciences: A Practitioners Guide” - *Short paper on project management by Economists, read it all today*
- 📖 Athey, S. and G. Imbens (ARE 2019): “Machine Learning Methods That Economists Should Know About” - *Well-written overview that introduces all the technical terms for machine learning, read it until 3rd day*
- 📖 Gentzkow, M., B. Kelly and M. Taddy (JEL 2019): “Text as Data” - *Well-written introduction to language processing, read it until last day*

How to use Python



Why Python?

- Interpreted, high-level, general-purpose programming language
- Can be object-oriented, imperative, functional and procedural
- Free (= no licenses)
- Large (= support and many packages)
- Centralized development
- Very good first language

Why Python?

- Interpreted, high-level, general-purpose programming language
- Can be object-oriented, imperative, functional and procedural
- Free (= no licenses)
- Large (= support and many packages)
- Centralized development
- Very good first language

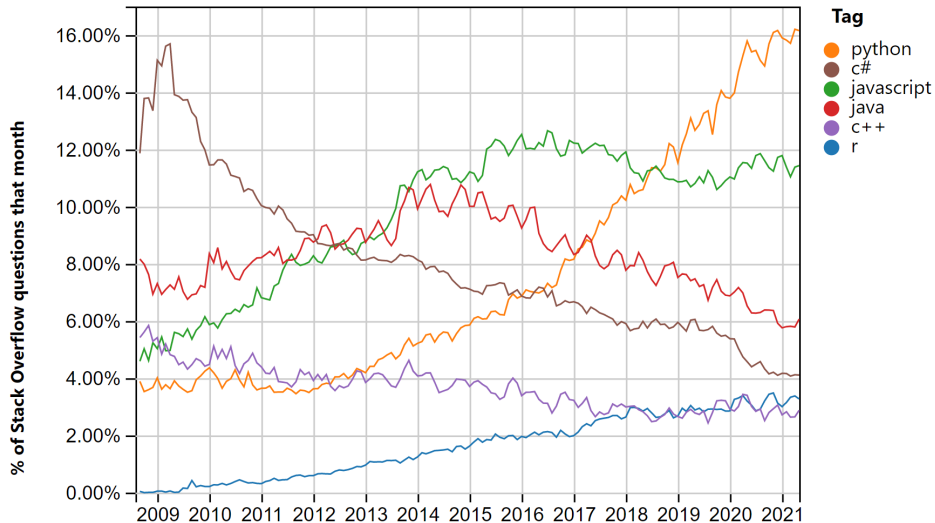
*There should be one— and preferably only one —obvious way to do it.
Although that way may not be obvious at first unless you're
Dutch.* (Tim Peters - The Zen of Python)

Credit where Credit is due

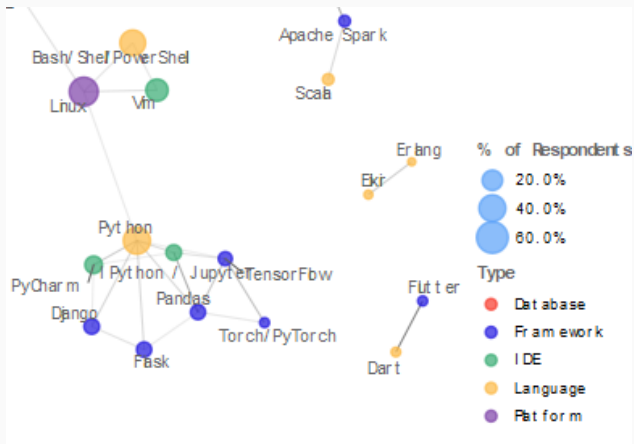
- Guido van Rossum
created Python in his Christmas holidays 1989 as
*"a descendant of ABC that would
appeal to Unix/C hackers. I chose
Python as a working title for the
project, being in a slightly irreverent
mood (and a big fan of Monty
Python's Flying Circus)."*
- Since 2019 5-member steering committee at
the Python Foundation heads the development
of Python



Python is popular and increasing in popularity



Python's local technology cluster



StackOverflow.com: ["Developer Survey Results 2019"](#)

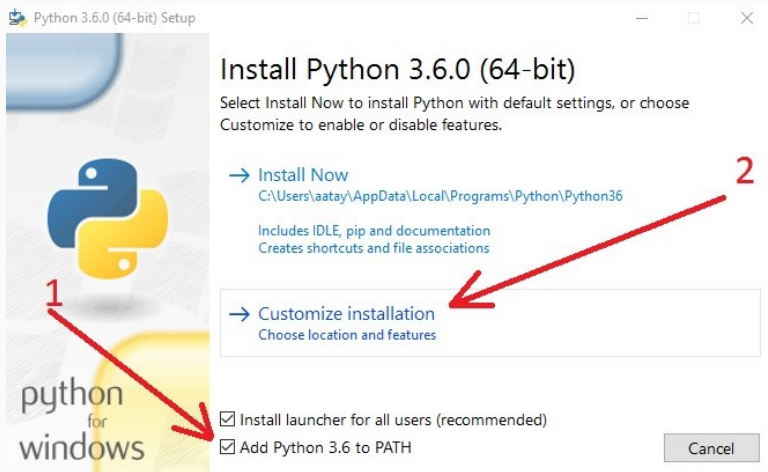
Why I discourage anaconda

- packages provided by anaconda need to be installed with `conda install` (they will ONLY be in the conda environment)
- packages tend to be outdated
- Overkill/Unnecessary software
- Jupyter and spyder run without anaconda as well
- Actually not *that* popular: 19% of Python installations via Anaconda¹


¹[Python Developers Survey 2020 Results](#)

Installing Python and pip

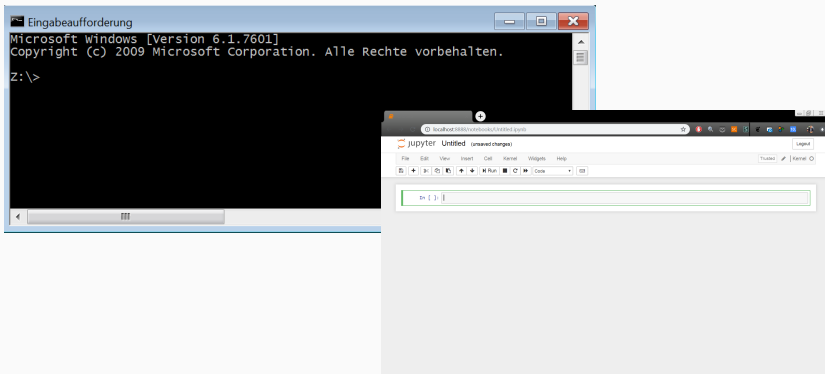
<https://www.python.org/downloads/>



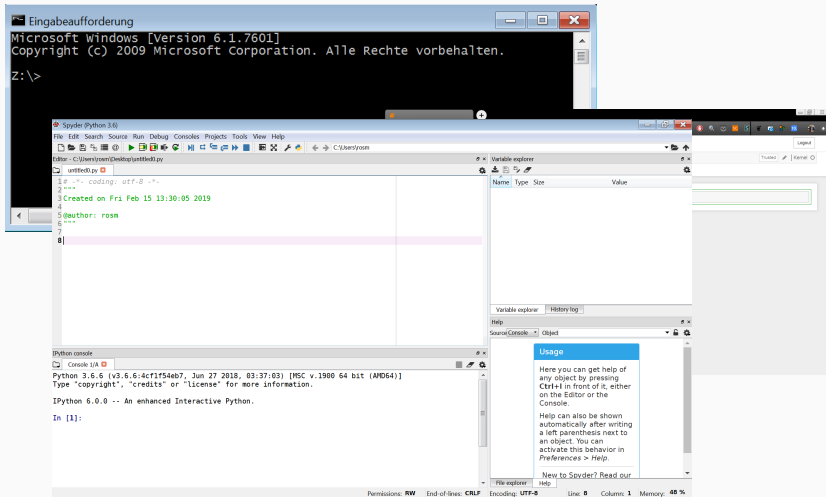
Different ways to use Python



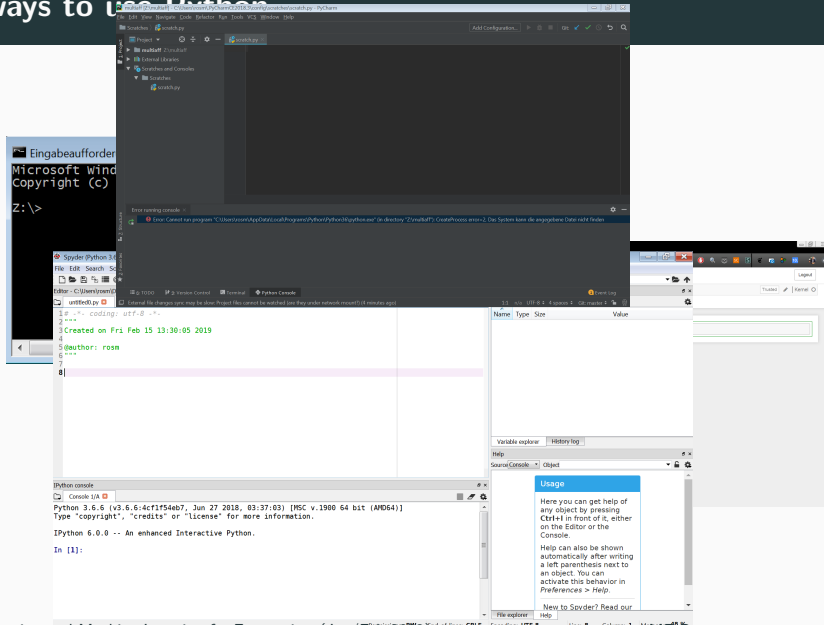
Different ways to use Python






Different ways to use Python



Different ways to use Python






Terminal/Console

- >_ Console uses DOS language () or shell and bash ( and )
- >_ Starts python environment, Jupyter, and executes scripts




>_ Console uses DOS language () or shell and bash ( and )

>_ Starts python environment, Jupyter, and executes scripts


>_ Install packages here:

 `python -m pip install pandas seaborn`
  `python3 -m pip install pandas seaborn`


>_ Shortcut (which is not platform-independent)

 `pip install pandas seaborn`
  `pip3 install pandas seaborn`

Jupyter Notebook on your computer


- Create a folder for this course and navigate there in your terminal (alternatively, open the "PowerShell" via context menu after +rightclick)

Jupyter Notebook on your computer

- Create a folder for this course and navigate there in your terminal (alternatively, open the "PowerShell" via context menu after +rightclick)
- Install the jupyter notebook if necessary

```
python3 -m pip install notebook
jupyter notebook
```
- Your browser will fire up (i.e., you started your own server)

Jupyter Notebook on your computer

- Create a folder for this course and navigate there in your terminal (alternatively, open the "PowerShell" via context menu after +rightclick)
- Install the jupyter notebook if necessary

```
python3 -m pip install notebook
jupyter notebook
```
- Your browser will fire up (i.e., you started your own server)
- Click on New in the upper right corner to start a new notebook

Notebooks will be saved in the folder where you invoked the jupyter server

- colab.research.google.com: requires Google account; stores notebooks in your Drive; integrates with GitHub; potentially older packages
- kaggle.com/code: requires Kaggle account; allows for R as well
- mybinder.org: requires GitHub account; builds from a GitHub repository

Recap some Python basics

What matters in Python?

- Indentation is key (convention: four spaces)
- Case-sensitive
- Variables must not start with numbers
- It's a language, *not* a program

Pandas



pandas: the library for data manipulation

- Documentation: <http://pandas.pydata.org/pandas-docs/stable/>

The screenshot shows the pandas documentation website. At the top, there is a navigation bar with the pandas logo and links for Home, What's New in 1.0.0, Getting started, User Guide, API reference, Development, and Release Notes. On the right side of the navigation bar are social media icons for GitHub and Twitter. Below the navigation bar is a search bar labeled "Search the docs ...". The main content area has the heading "pandas documentation" followed by the date "Date: Feb 05, 2020" and version "Version: 1.0.1". There are links for "Download documentation" (PDF Version | Zipped HTML) and "Useful links" (Binary Installers | Source Repository | Issues & Ideas | Q&A Support | Mailing List). A paragraph describes pandas as an open source, BSD-licensed library providing high-performance, easy-to-use data structures and data analysis tools for the Python programming language. Below this are four featured sections: "Getting started" with a person icon, "User guide" with an open book icon, "API reference" with a code icon, and "Developer guide" with a greater-than sign icon. Each section has a brief description and a button to access the content. At the bottom, there is a footer with the text "Python Programming and Machine Learning for Economists (Jan/Feb 2022)" and "ME Rose".

Home What's New in 1.0.0 Getting started User Guide API reference Development Release Notes

Search the docs ...


pandas documentation

Date: Feb 05, 2020 Version: 1.0.1

Download documentation: [PDF Version](#) | [Zipped HTML](#)

Useful links: [Binary Installers](#) | [Source Repository](#) | [Issues & Ideas](#) | [Q&A Support](#) | [Mailing List](#)


pandas is an open source, BSD-licensed library providing high-performance, easy-to-use data structures and data analysis tools for the Python programming language.



Getting started

New to pandas? Check out the getting started guides. They contain an introduction to pandas' main concepts and links to additional tutorials.


[To the getting started guides](#)



User guide


The user guide provides in-depth information on the key concepts of pandas with useful background information and explanation.

[To the user guide](#)



API reference

The reference guide contains a detailed description of the pandas



Developer guide

Saw a typo in the documentation? Want to improve existing

Python Programming and Machine Learning for Economists (Jan/Feb 2022) ME Rose

Let's start with a dataset on twins...

```
1 import pandas as pd
2
3 FNAME = "http://www.stat.ucla.edu/~rgould/datasets/twins.dat"
4
5 df = pd.read_csv(FNAME, sep='\t')
```

- Documentation at

<http://www.stat.ucla.edu/~rgould/datasets/twinsexplain.txt>

pandas functionality relevant for the course

- 10 minutes to pandas
- IO tools (text, CSV, HDF5, ...)
- Indexing and selecting data
- Reshaping and pivot tables
- Working with missing data
- Computational tools

Let's inspect our data

```
1 df.shape  # Dimensions
2 df.head() # First 5 lines (by default)
3 df.tail(7) # Last 7 lines
4 df.columns # List of variables
5 df.describe() # Summary statistics
```

1. How many observations do you have?
2. How many variables do you have?
3. Which variables are numeric?
4. What is the mean of variable "DEDUC1"?

Slicing the DataFrame

```
1  # Selecting columns
2  df["DEDUC1"]  # Column by column name
3  df[["AGE", "LHRWAGEH"]]  # Columns by list of column names
4  df.iloc[:, 5:7]  # Column range by column indices
5
6  # Selecting rows
7  df.loc[0]  # Row by index name (also accepts lists)
8  df.iloc[0]  # Row by row number (also accepts lists)
9
10 # Selecting values
11 df.loc[18, "AGE"]  # Name of row and column
12 df.iloc[18, 2]  # Index of row and column
```

Slicing the DataFrame

```
1 # Selecting columns
2 df["DEDUC1"] # Column by column name
3 df[["AGE", "LHRWAGEH"]] # Columns by list of column names
4 df.iloc[:, 5:7] # Column range by column indices
5
6 # Selecting rows
7 df.loc[0] # Row by index name (also accepts lists)
8 df.iloc[0] # Row by row number (also accepts lists)
9
10 # Selecting values
11 df.loc[18, "AGE"] # Name of row and column
12 df.iloc[18, 2] # Index of row and column
```

1. What is the 6th entry of the 5th column?
2. What is the 5th entry of column "DTEN"?
3. What is the last entry of column "LHRWAGEL"?

Understanding dtypes

```
1 df.info()
```

Understanding dtypes

```
1 df.info()
```

Pandas	Python	Purpose
object	unicode	Text
int64	int	Integers
float64	float	Floating numbers
bool	bool	True & False values
datetime64		Date and time values
timedelta[ns]		Differences between two datetimes
category		Finite list of text values

Changing dtypes

```
1 df["WHITEH"] = df["WHITEH"].astype(bool)
2 df["DMARRIED"] = df["DMARRIED"].astype("category")
3 df["LHRWAGEH"] = pd.to_numeric(df["LHRWAGEH"], errors="coerce")
```

Optimising dtypes

```
1 df.info(memory_usage=True)
```

Optimising dtypes

```
1 df.info(memory_usage=True)
```

```
1 bools = ['WHITEH', 'MALEH', 'WHITEL', 'MALEL']
2 df[bools] = df[bools].astype(bool)
3 df['DMARRIED'] = df['DMARRIED'].astype('int8')
4 df.info(memory_usage=True)
```

Boolean indexing

```
1 df[df["AGE"] > 20]    # One condition
2 df[(df["AGE"] > 20) & (df["WHITE"] == 1)] # Multiple conditions
3 df[~(df["AGE"] > 20)] # Tilde inverses boolean
4 values = (20, 21, 22, 23)
5 df[df["AGE"].isin(values)] # Select specific values
```

Boolean indexing

```
1 df[df["AGE"] > 20]    # One condition
2 df[(df["AGE"] > 20) & (df["WHITE"] == 1)] # Multiple conditions
3 df[~(df["AGE"] > 20)] # Tilde inverses boolean
4 values = (20, 21, 22, 23)
5 df[df["AGE"].isin(values)] # Select specific values
```

1. How many observations have "WHITE" equal to 0?
2. How many observations have "WHITE" equal to 1 and "DEDUC1" unequal to 0?
3. In how many rows do the values for "WHITE" and "WHITE" differ?
4. What is the mean age of twins whose L-sibling is a non-white male with either 12 or 14 years of education? (Use "WHITE", "MALE" and "EDUC",)

Aggregate data

```
1 df["WHITEH"].value_counts()  
2 pd.crosstab(df["WHITEH"], df["WHITEL"])
```

Aggregate data

```
1 df["WHITEH"].value_counts()
2 pd.crosstab(df["WHITEH"], df["WHITEL"])
```

1. What is the most common value in "EDUCL"?
2. What is the most common combination of "MALEH" and "MALEL"?

Manipulation

```
1  # Representation
2  df = df.sort_values(by='HRWAGEH')  # Sorting by column
3  df = df[sorted(df.columns)]  # Re-order columns alphabetically
4  # Work on columns
5  df = df.drop('AGESQ', axis=1)  # Drop a column
6  df['new'] = 9  # Add new column
7  df['AGETR'] = df['AGE']**3
8  df['combined'] = df['MALEH'] + df['EDUCH']
9  # Missing data
10 df["HRWAGEH_new"] = df["HRWAGEH"].fillna(0)  # Fill missings with 0
11 df = df.dropna(subset=["HRWAGEH"])  # Drop rows missing in "HRWAGEH"
```

Grouping

```
1 grouped = df.groupby(['MALEH'])
2 print(grouped['AGE'].mean())
3 print(grouped['EDUCH'].agg(['mean', 'sum']))
4 print(grouped[['EDUCH', 'AGE']].agg(['mean', 'std']))
```

Grouping

```
1 grouped = df.groupby(['MALEH'])
2 print(grouped['AGE'].mean())
3 print(grouped['EDUCH'].agg(['mean', 'sum']))
4 print(grouped[['EDUCH', 'AGE']].agg(['mean', 'std']))
```

- Full list at https://pandas.pydata.org/pandas-docs/stable/user_guide/groupby.html#aggregation
- What is the "AGE" variance for "MALEL" == 0 individuals?
 - What are the second and the third quartile of years of schooling for female L-siblings? (Use "EUDCL" and "MALEL" == 0)
 - What is the average "AGE" for twins where both siblings are female?

Creating DataFrames from other objects

Creating Pandas DataFrames from Python Lists and Dictionaries

	Dictionary		List																				
Row Oriented	<pre>sales = [{'account': 'Jones LLC', 'Jan': 150, 'Feb': 200, 'Mar': 140}, {'account': 'Alpha Co', 'Jan': 200, 'Feb': 210, 'Mar': 215}, {'account': 'Blue Inc', 'Jan': 50, 'Feb': 90, 'Mar': 95}] df = pd.DataFrame(sales)</pre>		<pre>sales = [('Jones LLC', 150, 200, 50), ('Alpha Co', 200, 210, 90), ('Blue Inc', 140, 215, 95)] labels = ['account', 'Jan', 'Feb', 'Mar'] df = pd.DataFrame.from_records(sales, columns=labels)</pre>																				
	default	<table border="1"><thead><tr><th></th><th>account</th><th>Jan</th><th>Feb</th><th>Mar</th></tr></thead><tbody><tr><td>0</td><td>Jones LLC</td><td>150</td><td>200</td><td>140</td></tr><tr><td>1</td><td>Alpha Co</td><td>200</td><td>210</td><td>215</td></tr><tr><td>2</td><td>Blue Inc</td><td>50</td><td>90</td><td>95</td></tr></tbody></table>		account	Jan	Feb	Mar	0	Jones LLC	150	200	140	1	Alpha Co	200	210	215	2	Blue Inc	50	90	95	from_records
	account	Jan	Feb	Mar																			
0	Jones LLC	150	200	140																			
1	Alpha Co	200	210	215																			
2	Blue Inc	50	90	95																			
Column Oriented	<pre>sales = {'account': ['Jones LLC', 'Alpha Co', 'Blue Inc'], 'Jan': [150, 200, 50], 'Feb': [200, 210, 90], 'Mar': [140, 215, 95]} df = pd.DataFrame.from_dict(sales)</pre>		<pre>sales = [['account', ['Jones LLC', 'Alpha Co', 'Blue Inc']], ['Jan', [150, 200, 50]], ['Feb', [200, 210, 90]], ['Mar', [140, 215, 95]]] df = pd.DataFrame.from_items(sales)</pre>																				
	from_dict		from_items																				

When using a dictionary, column order is not preserved.
Explicitly order them:
`df = df[['account', 'Jan', 'Feb', 'Mar']]`

Practical Business Python - pbpython.com

To become a Master...

🔖 10 minutes to pandas

📖 Wes McKinney: "Python for Data Analysis. Data Wrangling with Pandas, NumPy, and IPython", O'Reilly (2017)

📖 Fabio Nelli: "Python Data Analytics. Data Analysis and Science Using Pandas, matplotlib, and the Python Programming Language", Apress (2015)

Plotting w/ pandas (matplotlib), and w/ seaborn



Visualization with pandas

- Straightforward plotting as DataFrame methods for all kinds: barplots, areas, histograms, violin plots, timeseries, etc.:

<https://pandas.pydata.org/pandas-docs/stable/visualization.html>

- Has matplotlib under the hood - for aesthetics

```
import matplotlib.pyplot as plt
```

- Set global styles with `plt.style.use('<style>')` (list all styles with `plt.style.available`)

! Beware: Have DataFrame in correct format (long vs. wide)

Statistical plotting with seaborn

- [seaborn](#): wrapper for `matplotlib`, optimized for quick statistical plotting: Error bars, distributions, regressions, etc.
- Use seaborn's toy datasets using `.load_dataset()`
- 👉 If downloading example datasets via `.load_dataset()` doesn't work, get them from github.com/mwaskom/seaborn-data and store them in `~./seaborn-data/`

Seaborn's plotting philosophy

- Statistical relation between numeric values?
 - ➔ `relplot()` for Scatter and Line (→ [Documentation](#))
- Categorical data?
 - ➔ `catplot()` for Scatter-like (Swarm and Strip), Distributions (Box, Violin, Boxen) and Estimations (Point, Bar, Count) (→ [Documentation](#))
- Linear relationships?
 - ➔ `regplot()` (→ [Documentation](#))

Pandas plotting vs. seaborn

- In Jupyter, remember to write and execute `%matplotlib inline` in first cell to show figures
- Use pandas when you do the aggregations yourself
- Use seaborn when you use raw data – seaborn will aggregate itself

Color maps

List of named colors

To become a Master...

 Fabio Nelli: "Python Data Analytics. Data Analysis and Science Using Pandas, matplotlib, and the Python Programming Language", Apress (2015)

 matplotlib [Tutorials](#)

 seaborn [User guide and tutorial](#)